

# Profonda Commedia

Deep learning project on Natural Language Generation of canticas from Divina Commedia using Dante's style

(<https://github.com/dropino/ProfondaCommedia>)

Nicholas Carroll

[nicholas.carroll@studio.unibo.it](mailto:nicholas.carroll@studio.unibo.it)

Simone Vagnoni

[simone.vagnoni2@studio.unibo.it](mailto:simone.vagnoni2@studio.unibo.it)

## Abstract

In this work we developed two models able to generate texts similar to Dante's Divine Comedy, with particular attention given to making sure that verses are correctly syllabified.

The technique used to accomplish these systems is the Transformer, a neural network architecture that is widely used in Natural language processing, and also that gives the best results, according to the literature on the subject.

Two different architectures have been developed: the first is composed of a **single Generator model trained on pre-syllabified data** to directly generate syllabified verses. The second one is composed of a **Generator trained with non syllabified data** which generates verses and a **Syllabifier trained with syllabified data** to demonstrate that the generated verses are correct.

We trained the network on a dataset consisting in the Divine Comedy (both divided in syllables and not) and several experiments were conducted to tune the hyperparameters to get the best possible results.

# 1. Data

To train our networks we used two different versions of the divine comedy. One is a plain, non syllabified version of the divine comedy, while the other was the syllabified version created by professor Asperti (<https://github.com/asperti/Dante>) that we used to train both the Generator in the all-in-one architecture that generates syllabified text, and the Syllabifier that in the other architecture syllabifies the verses generated by the unsyllabified Generator.

The dataset is quite small as it consists of only 14 233 verses. We had also planned to use other works of Dante or other authors to increase the amount of available data but their different structure and style proved troublesome when trying to generate text in the style of the divine comedy.

The dataset was divided in training set, validation set and dataset. Since there was not much data we divided the verses in 95% for the training set, 3% validation set, 2% testing set.

## 1.1 Cleaning the data

The first step in using the text was to clean the data by removing the verse numbers, the titles of each canto (for example: "inferno Canto i") and any special characters such as •, —, -, (, ), « and ». Punctuation was also removed except for question marks and exclamation marks.

## 1.2 Syllabification

Initially we worked with a plain text version of the divine comedy and wrote an algorithm to syllabify it. Results were acceptable but we encountered some problems with the synalepha given its dependence on the context.

Synalephas are a metric figure in which, when counting syllables of a verse, two syllables are merged. If the figure is not respected the verses end up having the wrong number of syllables.

Given these difficulties we decided to use an already hyphenated version of the divine comedy made by professor Asperti. Training the networks with this version generated text that better followed the hendecasyllabic structure of the divine comedy and most importantly had less problems with synalepha.

## 1.3 Tokenization and creating the dataset

A model can't be trained directly with text, it has to be converted to a numeric representation. This is usually done by transforming the text in a sequence of numeric token IDs.

At first we tokenized the text at a syllabic level using the syllabified dataset we had but it proved to be less than ideal as it led to numerous repetitions within verses, generating examples such as “e così così così così così così”. To try and improve the generation we evaluated two main strategies, tokenizing the texts letter by letter and dividing the verses to create a subword vocabulary.

The main advantage of a subword tokenizer is that it interpolates between word-based and character-based tokenization. Common words get a slot in the vocabulary, but the tokenizer can fall back to word pieces and individual characters for unknown words. To do so we used the **BertTokenizer**, implemented as a class by TensorFlow with a higher level interface. It includes BERT's token splitting algorithm and a WordPieceTokenizer. It takes **sentences** as input and returns **token-IDs**.

After numerous tests to evaluate the best solution we ended up using letter tokenization in both generators and the Bert Tokenizer in the syllabifier.

We used a dictionary to map all the letters and symbols to an ID, ending up with 42 tokens. In this amount we also included 4 special tokens:

- “[SEP]” for each space
- and inserted at the beginning of each verse the token “[START]”.
- To make all verses the same lengths, we used the special character “[PAD]” to bring them up to the length of 100 tokens.
- At the end of each verse we appended the symbol “[STOP]”
- “[UNK]” for unknown tokens that are not present in the dictionary
- “[EOS]” for the end of a sentence.

For the syllabizer we created two different dictionaries as we treated the problem as a translation problem, transforming the non-syllabized text to a syllabized version. Each of these dictionaries has the same special tokens as the letter tokenizer except for “[EOS]” and “[SEP]”.

The syllabized text vocabulary is composed of 796 tokens while the non-syllabized text vocabulary is composed of 1738 tokens.

## 2. Model

Reading the literature about this subject, we learned that the best results for this kind of problems were given by the Transformer model, instead of the traditional RNN (Long-short term memory LSTM or GRU). One of the main reasons for this choice is that RNN processes one symbol at time and generates a sequence of hidden states  $h_t$  as a function of the previous hidden state  $h_{t-1}$  and the input for position  $t$ . This sequential way to operate **precludes parallelization** and it is **critical with long sequences**: the unrolled representation of a RNN results in a very deep neural network, leading to problems like vanishing and/or exploding gradients that are only partially corrected by the gated units. Instead the **Transformer** architecture, besides the parallelism with which it deals with the input data, can more easily catch the long term dependencies between complex syllabic structures.

### 2.1 Our architectures

As previously mentioned our work consists in two distinct architectures which we have called the syllabifier-generator and the all-in-one generator for convenience.

#### 2.1.1 Syllabifier and non syllabified generator

So the model chosen is an architecture that ensembles **two different Transformers, a Syllabifier and a Generator**, that respectively translate the input in a syllabified output, and create verses with the style of the input. In particular the Generator, trained with the un-syllabified Divine Comedy, creates verses that syllabified from the Syllabifier, reach the objective of the task, that can be compared with the one of the **Generator** of the verses starting from the Divine Comedy Syllabified.

The first one, the Syllabifier, takes a file of the verses mapped (un - syllabified as) an input, that after being tokenized and **encoded**, is given to the transformer to train to associate the two texts (with and without the “|” symbol to separate syllables), and predict the most accurate “translation”, that after being **decoded**, gives the syllabified version of the text.

The other model, the Generator, takes as input the syllabified Divine Comedy, which is tokenized and the data is fed to the **encoder** that maps the input

sequence of symbol representations to a sequence of continuous representations. Given that continuous one, the **decoder** then generates an output sequence of symbols one element at a time. At each step the model is auto-regressive, consuming the previously generated symbols as additional input when generating the next. The Transformer follows this overall architecture using stacked self-attention and point-wise, fully connected layers for both the encoder and decoder

### 2.1.2 All-in-one generator

The architecture for the all-in-one generator is practically identical to the Generator in our other Generator-Syllabifier model. The only differences lie in the pre-processing of the data and a slightly different implementation.

## 2.2 Positional encoding and embedding

Since this pattern contains no recurrence or convolution, positional encoding is added to provide the pattern with some information about the relative position of the words in the sentence. The positional encoding vector is added to the embedding vector. Embeds represent a token in a d-dimensional space where tokens with similar meaning will be closer to each other. But embeddings don't encode the relative position of words in a sentence. Thus, after adding the positional coding, the words will be closer to each other based on the similarity of their meaning and their position in the sentence, in the d-dimensional space. The formula for calculating positional encoding is as follows:

$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{model}})$$

$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{model}})$$

Beside the I/O sequences, a boolean *mask* is fed to the model in order to ignore padding symbols. All the pad tokens in the batch of sequence are masked. It ensures that the model does not treat padding as the input. The mask indicates where pad value 0 is present: it outputs a 1 at those locations, and a 0 otherwise.


## 2.3 Multi-head, scaled dot product attention

Multi-head attention consists of four parts:

- Linear layers and split into heads.
- Scaled dot-product attention.
- Concatenation of heads.
- Final linear layer.


Each multi-head attention block gets three inputs; Q (query), K (key), V (value). These are put through linear (Dense) layers and split up into multiple heads.

The scaled-dot-product-attention defined above is applied to each head (broadcasted for efficiency). An appropriate mask must be used in the attention step. The attention output for each head is then concatenated and put



through a final Dense layer.

### Scaled Dot-Product Attention

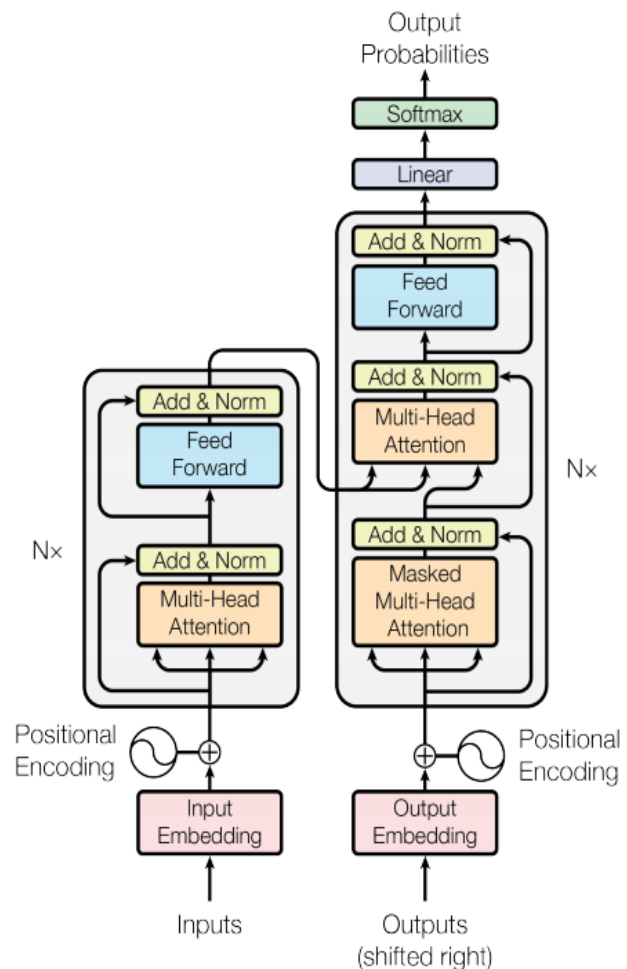


Instead of one single attention head, Q, K, and V are split into multiple heads because it allows the model to jointly attend to information from different representation subspaces at different positions. After the split each head has a reduced dimensionality, so the total computation cost is the same as a single head attention with full dimensionality.

## 2.4 Transformer

A transformer is a deep learning model that adopts the mechanism of attention, differentially weighing the significance of each part of the input data. It is used primarily in the field of natural language processing (NLP) and in computer vision (CV).

Like recurrent neural networks (RNNs), transformers are designed to handle sequential input data, such as natural language, for tasks such as translation and text summarization. However, unlike RNNs, transformers do not necessarily process the data in order. Rather, the attention mechanism provides context for any position in the input sequence. For example, if the input data is a natural language sentence, the transformer does not need to process the beginning of the sentence before the end. Rather, it identifies the context that confers meaning to each word in the sentence. This feature allows for more parallelization than RNNs and therefore reduces training times.



## 2.5 Encoder

Each encoder consists of two major components: a self-attention mechanism and a feed-forward neural network. The self-attention mechanism accepts input encodings from the previous encoder and weighs their relevance to each other to generate output encodings. The feed-forward neural network further processes each output encoding individually. These output encodings are then passed to the next encoder as its input, as well as to the decoders.

The first encoder takes positional information and embeddings of the input sequence as its input, rather than encodings. The positional information is necessary for the transformer to make use of the order of the sequence, because no other part of the transformer makes use of this.

## 2.6 Decoder

Each decoder consists of three major components: a self-attention mechanism, an attention mechanism over the encodings, and a feed-forward neural network. The decoder functions in a similar fashion to the encoder, but an additional attention mechanism is inserted which instead draws relevant information from the encodings generated by the encoders.

Like the first encoder, the first decoder takes positional information and embeddings of the output sequence as its input, rather than encodings. The transformer must not use the current or future output to predict an output, so the output sequence must be partially masked to prevent this reverse information flow. The last decoder is followed by a final linear transformation and softmax layer, to produce the output probabilities over the vocabulary.



## 3. Generator

Our Generator is based on the Transformer, an **auto-regressive** model, meaning that it will predict the next token based on the past sequence already generated (decoder input).

In the all-in-one architecture, in order to generate a sample of text, we fed the decoder with an unseen starting sequence from the test set and then concatenated gradually to the input the token sampled according to the model prediction.

### 3.1 Top-K sampling

In order to choose the best syllables to append to the sequence, but not always the same, we used *Top-K* sampling, in which the  $K$  most likely next words are filtered and the probability mass is redistributed among only those  $K$  next words. GPT2 adopted this sampling scheme, which was one of the reasons for its success in story generation.

The most problematic aspect of TopK search is the worsening of results if it maintains the same  $K$  for every token to be generated. The best implementation of this would be to dynamically adapt the value of  $K$ , decreasing it toward the end of the sentence where the rhythmic scheme must be preserved. A trick is to make the distribution of chosen words sharper (increasing the likelihood of high probability words and decreasing the likelihood of low probability words) by lowering the so-called **temperature** of the softmax.

Setting the parameters  $k$  and  $t$  (temperature) to 1 the result is a **greedy** algorithm that chooses the most probable word, but due to its simplicity, could often lead to loops of sequences of same or similar words.

## 4. Syllabifier

The structure of the syllabifier is also a transformer with the same modalities of the previous one that using the BertTokenizer that embeds the verses in encodings, translates the unsyllabed verses into the syllabed ones. This is made training the network to associate the couples of verses (syll. and unsyll.).

In particular the steps that the transformer made to infer the next tokens are:

- Encode the input sentence using the Unsyllabified tokenizer (tokenizers.unsyll). This is the encoder input.
- The decoder input is initialized to the [START] token.
- Calculate the padding masks and the look ahead masks.
- The decoder then outputs the predictions by looking at the encoder output and its own output (self-attention).
- Concatenate the predicted token to the decoder input and pass it to the decoder.
- In this approach, the decoder predicts the next token based on the previous tokens it predicted.

# 5. Training

The transformers within the generator and the syllabifier were both trained with Google **Colab**.

Every 5 epochs we validated the models using the validation dataset and averaging the results. At prefixed epochs the models are used to generate a sample of text, in order to keep track of the performance at different times during training.

We kept track of train accuracy, train loss, validation accuracy and validation loss during the different training sessions in a consistent and ordered way using the **Weights and Biases** API, that allowed us to log and compare metrics, hardware statistics and generated data in a web application.

All networks were trained with batches of 8 tercets for a total of 24 verses at a time.

## 5.1 Generator hyperparameters

Like in classical Natural Language processing tasks, the loss function is the sparse **categorical cross entropy** between the predicted tokens and the padded masked target sequence.

We used the **Adam** optimizer with  $\beta_1 = 0.9$ ,  $\beta_2 = 0.98$  and **epsilon** =  $10e-9$  and a fixed learning rate of  $2e-4$ , tuned during several experiments.

The other hyperparameters have been set as follows:

- number of encoder/decoder layers = 4
- dimension of Embedding and attention layers = 128\*
- dimension of Feed Forward inner layers = 256
- dropout rate = 0.1
- number of attention heads = 4
- epochs = 50

The model, written in python, was developed using the **Keras** framework on top of **Tensorflow** 2.0.

## 5.2 Syllabification hyperparameters

The hyperparameter chosen for the generators are:

**-loss function: Sparse categorical cross-entropy**

We used the **Adam** optimizer with  $\beta_1 = 0.9$ ,  $\beta_2 = 0.98$  and **epsilon** =  $10e-9$  and a learning rate that varies based on this formula (<https://arxiv.org/abs/1706.03762>)

$$lrate = d_{model}^{-0.5} * \min(step\_num^{-0.5}, step\_num \cdot warmup\_steps^{-1.5})$$

**-number of encoder/decoder layers = 4**

**-dimension of Embedding and attention layers = 128\***

**-dimension of Feed Forward inner layers = 512**

**-dropout rate = 0.1**

**-number of attention heads = 8**

**-epochs: 20**

With these setting we have a syllabization correctness of 97%.

## 6. Results

After several attempts, changing the hyperparameters of the consolidated base architecture (dropout in [0.1, 0.3, 0.5], learning rate in [1e-3, 2e-4], attention heads in [2, 3, 4]) we obtained the best results, in terms of the text evaluation metrics, with the reported configuration. The loss and the accuracy were respectively **0.67** and **0.84**.

But the hyperparameters that we found most influential in terms of the structure of the hendecasyllables generated were the evaluating parameters of the Top-k function: the number **k** and the **temperature t**.

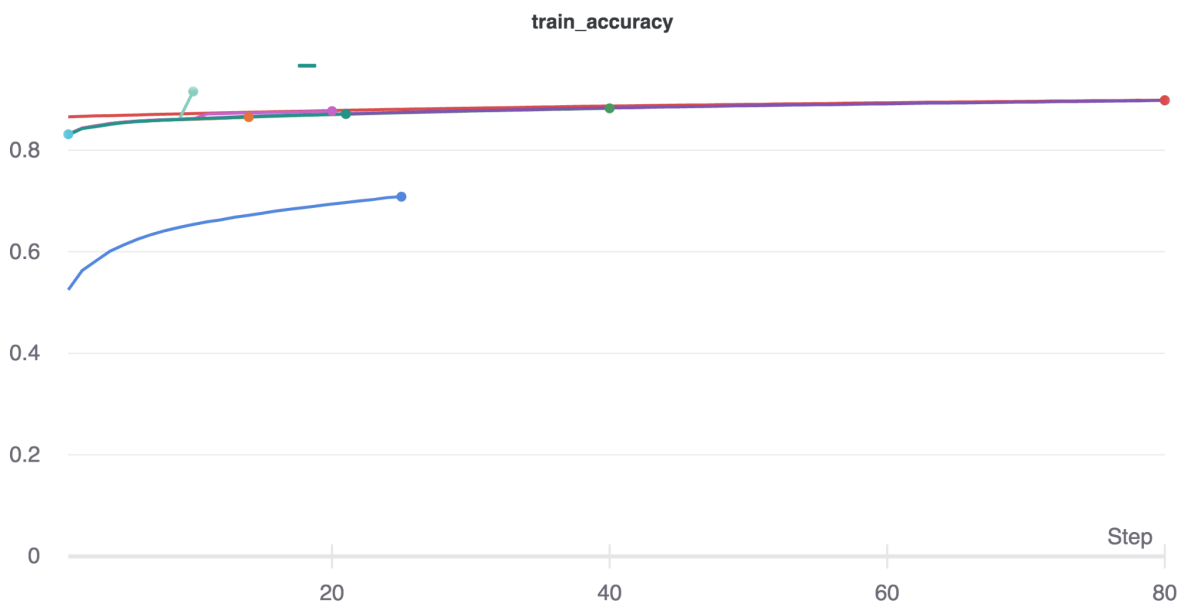
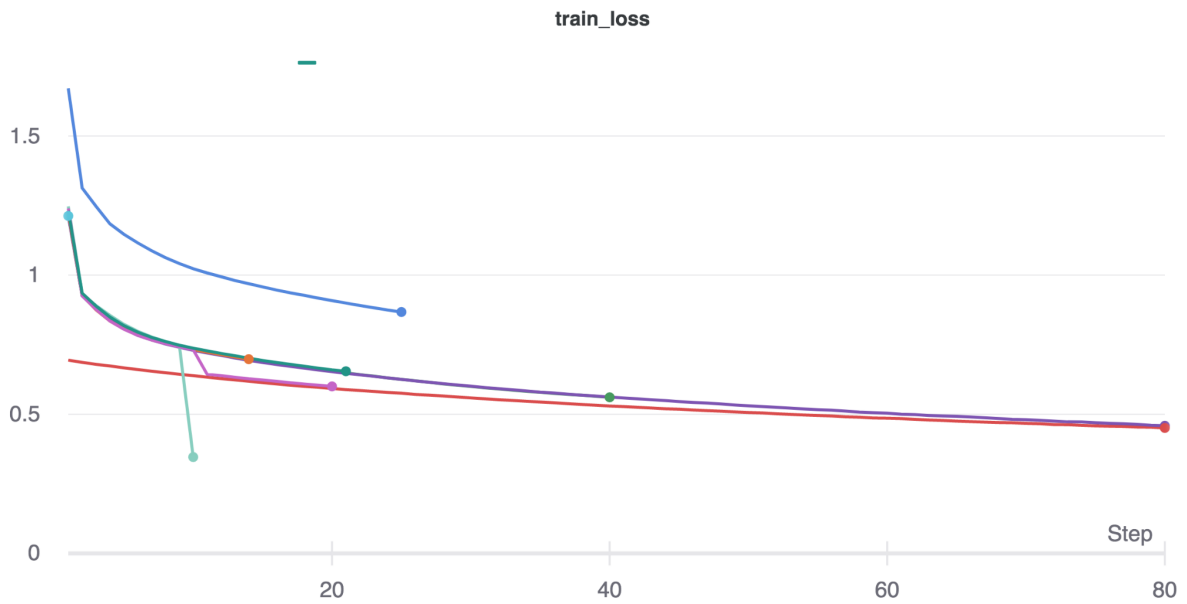
Ranging the k from 1 to 15 and the temperature from 0 to 1, we found the best results with **k=3** and **t=0.9**, because with higher k or lower t, the results where a monotonous repetition, often looping of the same words and syllables or the tercets were formed by the wrong number of verses.

Here following the text generated by the all-in-one architecture:


<p>or   que  sti   che   da  l' in  fi  ma   la  cu  na </p> <p>de  l' u  ni  ver  so in  fin   qui   ha   ve  du  te </p> <p>le   vi  te   spi  ri  ta  li ad   u  na ad   u  na </p>	<p>ci  s'   sen  do   cie  se el  l' on  ta  ra  ro </p> <p>che'   pri  scon   cial   co  ra  ra  ro   so  se   pia  sta </p> <p>es  si   la   son   son  do   so  mo  ra al   so  ro </p>
<p>che  s' o  ria   che   cia   so  ma  sto   se  sce </p> <p>co  sci   se  scen  ter   la   sol  to   che   la  sto </p> <p>e   la   so  sta   co  scia  sto   suo   su  ra  ta </p>	<p>se   ciò   che  ch'i'   lo   con  tro  ren  che   pren  to </p> <p>e   la   che   son   la  ch' io   che   che   si   per  to </p> <p>che   si  l' u  ma   sua   cor   piu   la   lo   pre  so </p>
<p>che   chi  l' io   la   son  to   co s'   sol  ta </p> <p>e   la   si  ri  l' o  so a  spa  no all  cia   si  ra  so </p> <p>si'   che   con  te   se  l' a  ve  ta e  stre   spe  ro </p>	<p>che   con  tre'n   cia  ri  tra   si   se   se  sce </p> <p>e   con   co  re  so   chi   chi   sol  la   lo  ri  sta </p> <p>e  l' a  ma   su  tra   li   la  tro   ci   so  stro </p>

e  l' al  li   chi   per   per  so  so   che   per  se	e  l' al  ta   per  ta   cia  stra  l' a   pe  sto
che   son   cio  l' al  ta  ra  star   sua   chia  re	si   che   si  stra  stre  l' al  tra   pe  ne   son   chio
el  ch' al  tra   pri  to   cor  sta   so  ra  sto   cor  to	si  ch' a  ri  tri  l' al  li  to   sia  to   par  ta
co  sce   son  tra   so  sto  l' a  me   la   con   pri	se  ve  ste  sto   con   so  ra e   co  ro  sce   cor   co  so
e   se  ni   son   son  l' al  tra  se   so  la	se   suo  ra   se  mi  ch' al   co  stra   pe  ro
e  ch' a  mi  tra   so  me   so  li   co  ra   chi   pier  to	si  l' ac  co   la   ci  ch' a  ra  ch' a   so  sto
sol  tra  l' al   pe  ra   si  l' ol   per  co a  ra  to	co s'   la   si   la   suon   se   la  tre  te   pie  ta
e  ch' al  tar  si   so  sto a   cor   ser  si  to	e   la   chi  l' a  vi  ra  l' ac  ci   sol   so  se
co s'   so  le al  tri  sta   che   cor   sen  to	e   la  mi  to e   la   son   pri  sta  ta   con  te
el  lo  ro   co  men  te   la   che  l' io   si   pa  ta	lo   sia   con  tra  se  l' al  li al   si   che'l   se  te
e   lo  ch' e  si  sto   sol  la   so  se   per  no	co  ra   lo  stro  sce  ra  ste  ra a  ra   cor  ta
che   che   li   che   suon  di  tre  re  sto   co  ste	co  ra   se  ra   su   li   che   co  sche   se  sca
che   la   se   co s'   la   cia   la  ch' al  ten  te  sta	che  l' u  si   che   so  sta   ci   so  stra   spre  sto
co  sto   la   che   le   che   si   so  ra  sce   spe  scio	e   con  do al  te   suo  ra   so  se   cian  de
per i  mon   che   co  me   so  mo  ra  sto   mi   mor  te	si   se  ra   la  ch' io   se   suo  ra   co  re
la   la'l   per   so  se   ser  che   che   che   spa  no	di  stran  di   che   li   co  me  s'  l' in  ch' o  sten  to
ch' io   che   con  dor  sta  sto  l' a  se a  stra  ra	e   la   la   lo al   pri  ta  stra   la   so  ste
per  tor   che   che   se   la   con   se  si   sol  le	e   la   la   lo   so  ren  ta   ser  se   so  stra  so
co  me   che   che'l   con   si   la  re a  me   con  de	co s'   la   la   che   se   la   sun  ta   chio  se

<p>che a  se e   suo   con  l' a  ste   con  che   co  re </p> <p>ch' ion  chi   si   la'l   piu   per   so  so   si   ca  spo </p> <p>e   che  l' al   sol  ch' al  la   prar  ch' o  ri  to </p> <p>che   so  ste   so  sci e   son   so  mi  l' al  li   spi  ra </p> <p>so  ste  ch' a   suo   la   che   che'n   so  re </p> <p>co s'   co  stra  re   cor  ci   le   son   son  ta </p> <p>che   se in  di  ch' io   si  l' al  tri   son  ta  ra </p> <p>cial   piu   suo   si   chi   la   la  stre   se  ra   per  se </p> <p>con  l' al  tri  sto   co  ran  to   chia   con   co  ste </p> <p>si  ch' a  stra   cio   so  ro   si   se  re  ra   so  na </p> <p>per   sen   la   si   si   la   so  me   son   si   par  si </p> <p>che   si  sia  to   se   sen   suo   son  to   so  sto </p> <p>ch' a  mi   cor  ta  ro a  l' a   suo  ra  sto   car  so </p> <p>co s'   che   co  stre   la   son   so  ra  so   pa  sto </p>	<p>si   sol   si   co  me   co  stran  tra   che   sen  to </p> <p>che al  ch' a  ster   se  ste   se   con   so  to </p> <p>per   la   che   che   so  stra  l' a   so  sco  so </p> <p>che  l' a  si   piu  l' o  si  l' al  tri  se   piu'   coll  li </p> <p>chi   si   si   con   co  me   che e   su  si   sie  na </p> <p>di  l' al  ta   la  scia   la  l' a   pier  tro   sper  so </p> <p>e  l' o  spri   che   la   pra  ra a  se  ra  sten  ta </p> <p>co s'  l' a  ri  sta  ra   sia  se   la   con  do </p> <p>coi  me   la   suo  na   con   se  re   se  ne  na </p> <p>se   che   ci   se  ra  sta  ne   le   la   sun  to </p> <p>e   cor  so   co s'   si   che   cial  ta a  mor  te </p> <p>la  l' al  lo   piu  che   si   su  tor  ta  ra </p> <p>e   che   lar  si  te   se  ra   chia   se   si   sen  to </p> <p>se  s'   la   so  ra   la   su  ra  te   la   par  sa </p> <p>la   sor   si  mi  to a   co  ma   lo   con   sor  ca </p>
--	---







Here reported in the figures the obtained loss and accuracy of the train and validation set of the all-in-one generator, with respect of the epochs.

Here following the text generated by the Generator-Syllabifier architecture:

nel  qual  non  si  dee  cre  der  che  s ' in  vii	che  la  suo  ra  son  tran  to  sen  che  la  sua
---	--

<p> per  cre  a  tu  ra  l'  oc  chio  tan  to  chia  ro</p> <p> e  io  ch'  al  fi  ne  di  tut  t'  i  di  sii</p> <p> che  per  con  do  di  co  me  che  la  man  to</p> <p> che  sen  che  l'  u  na  sol  tro  sua  re  na</p> <p> po  re  l'  an  co  al  con  che  per  con  de  na</p> <p> se  per  la  so  re  sor  dal  suo  e  stra  suo</p> <p> e  con  che  per  che  se  se  son  tro  se  glio</p> <p> ser  sen  to  a  co  sta  se  per  san  con  de</p> <p> che  la  man  to  e  la  suo  al  tro  sen  ta  ta</p> <p> ch'  e  nis  se  chi  se  ra  che  la  son  sor  te</p> <p> e  sar  ra  te  suo  co  me  com'  a  sua  sua</p> <p> che  la  sor  che  l'  an  to  e  scen  do  sor  ta</p> <p> con  l'  e  ra  con  sen  do  se  la  mar  te</p> <p> co  ri  che  sa  sen  to  e  che  su  ra  suo  ra</p> <p> che  per  se  re  son  per  le  la  su  ra  na</p> <p> si  son  con  do  co  men  to  a  la  son  par  ta</p> <p> co  si  son  suon  se  san  do  e  per  che  per  rar  ra</p> <p> e  sa  re  con  sem  pen  se  suon  de  an  ta  sen  te</p> <p> e  che  l'  u  ne  so  le  dal  cos  sa  l'  al  tro</p>	<p> e  con  se  la  man  suon  de  ra  la  su  ra</p> <p> e  la  ser  ra  l'  ar  ran  te  chi  per  sar  ra</p> <p> e  ze  che  l'  an  che  se  gua  suo  le  per  te</p> <p> che  par  sen  te  son  ser  se  na  se  re</p> <p> e  co  me  re  la  suo  son  la  so  rem  ta  rar  se</p> <p> che  se  re  la  sen  son  che  se  l'  er  sen  suo  to</p> <p> e  se  pa  rem  po  re  che  le  se  la  suon  te</p> <p> si  son  sen  l'  a  rar  di  co  men  ten  den  de</p> <p> e  che  son  tro  a  son  de  le  sol  ten  par  te</p> <p> che  le  sua  con  se  so  re  sua  suo  na</p> <p> ch'  al  suo  che  son  de  con  te  car  ta  ta</p> <p> con  con  se  mer  son  se  san  la  pren  to</p> <p> el  per  sen  che  son  do  se  l'  et  ta  ser  let  ta</p> <p> con  son  sar  sua  le  le  sua  ra  se  ra  ra</p> <p> de  per  le  ter  sen  to  e  per  ch'  io  sor  te</p> <p> ch'  el  suo  son  che  sen  de  l'  al  tror  si  son  ta</p> <p> e  l'  an  tra  suon  der  le  par  to  son  ter  ne</p> <p> per  se  co  si  di  sen  tra  la  ser  car  te</p> <p> e  la  par  te  ra  con  san  ta  e  sal  cas  so</p> <p> con  de  sen  te  sen  per  en  te  lan  te  cor  so</p>
---	---

<p> con  ser  re  san  do e  sar  che  la  su  ra</p> <p> che  co  man  tra  con  che  pon  sen  so  le</p> <p> e  che  par  te  su  l ' en  con  su  se  son  sen  te</p> <p> con  sua  che  l ' es  si  cos  se  pre  stre  la</p> <p> che  la  suo  an  te  suon  suon  suon  al  tro</p> <p> di  so  ren  te  che  le  son  do a  suon  le  sem  mo</p> <p> con  che  suar  che  par  ta  sen  te  pe  re</p> <p> e  se  suon  sen  dor  che  di  su  se  ra  sen  na</p> <p> che  sor  te  che  su  se  se  la  sor  to a  sol  tra</p> <p> con  del  del  son  con  cos  sa  per  per  ta  ta</p> <p> che  l ' al  man  do e  l ' an  dan  son  son  so  te</p> <p> e  l ' an  tra  co  me  la  suo  al  tre  suo  ser  ro</p> <p> con  son  to  che  par  la  mor  sor  so  ra</p> <p> che  suon  con  to  se  sol  son  son  suan  ta</p> <p> di  com '  an  ten  da  chia  che  la  par  son  ta</p> <p> che  la  suo  suol  che  san  te  ser  re  pen  ta</p>	<p> per  che  che  por  te  se  suon  do  so  re</p> <p> e  san  tro an  to  con  si  can  do  ser  len  te</p> <p> e  sen  co  sto  al  te  sor  di  per  sua</p> <p> co  me  con  se  ma  la  ser  che  par  l ' al  tro</p> <p> se  per  la  sur  lar  la  son  por  to al  per  sa</p> <p> e  l ' ar  ra  la  par  che  le  suo  se  men  te</p> <p> e  san  de  ro  sua  son  son  suon  con  per  to</p> <p> ch ' io  che  le  tu  son  co  re  suon  so  tro</p> <p> son  sen  l ' al  la  sen  to an  co a  sua  san  te</p> <p> che  son  cam  bra  sua  ten  te  che  le  se  igno</p> <p> se  san  do al  suo  ra  par  con  di  ser  per  co</p> <p> che  per  l ' al  tro  son  to e  la  sem  sem  pe</p> <p> e  sar  sen  co  me  se  la  ca  ren  che  so  ra</p> <p> co  me  per  la  man  tra  la  par  l ' ar  ta  ta</p> <p> e  san  te  l ' on  da  con  sen  to a  per  se  ra</p>
--	--

## 6.1 Text evaluation metrics

To evaluate the quality of the generated cantica we used mainly two metrics ranging in the interval  $[0..1]$ , these however cannot replace human-based judgement.

- **Terces structureness**: the ratio between the number of well formed terces intended as 3 sentences separated by an empty new line and the expected number based on the total lines produced.

- **Hendecasyllabicness**: average compliance of each verse to the hendecasyllable pattern.

These are the results we obtained with the All-in-one architecture:

**Terces structureness: 0.89**

**Hendecasyllabicness: 0.964**

These are the results for the Generator-Syllabifier architecture:

**Terces structureness: 0.86**

**Hendecasyllabicness: 0.963**

## 7. Conclusions

Both models give very similar and satisfactory results with our metrics but as previously mentioned human based judgement is the best metric in this case. In particular we believe that the Syllabifier-Generator based model is the best one in this case.

## 8. Bibliography and Sitography

1. Vaswani A. et al, Attention Is All You Need, 2017  
<https://arxiv.org/abs/1706.03762>
2. Asperti A., del Bianco S., Syllabification of the Divine Comedy, 2021  
<https://arxiv.org/abs/2010.13515>
3. Minh-Thang L., et al. Effective Approaches to Attention-based Neural Machine Translation, 2015 <https://arxiv.org/abs/1508.04025v5>
4. [https://www.tensorflow.org/text/tutorials/transformer#create\\_the\\_transformer](https://www.tensorflow.org/text/tutorials/transformer#create_the_transformer);
5. [https://www.tensorflow.org/text/tutorials/text\\_generation](https://www.tensorflow.org/text/tutorials/text_generation);
6. [https://www.tensorflow.org/text/tutorials/nmt\\_with\\_attention](https://www.tensorflow.org/text/tutorials/nmt_with_attention);
7. [https://www.tensorflow.org/text/guide/subwords\\_tokenizer#build\\_the\\_tokenizer](https://www.tensorflow.org/text/guide/subwords_tokenizer#build_the_tokenizer)
8. [https://github.com/DanieleVeri/deep\\_comedy/blob/master/report.pdf](https://github.com/DanieleVeri/deep_comedy/blob/master/report.pdf)
9. <https://machinelearningmastery.com/adam-optimization-algorithm-for-deep-learning/>
10. <https://blog.ml.cmu.edu/2020/03/20/are-sixteen-heads-really-better-than-one/>